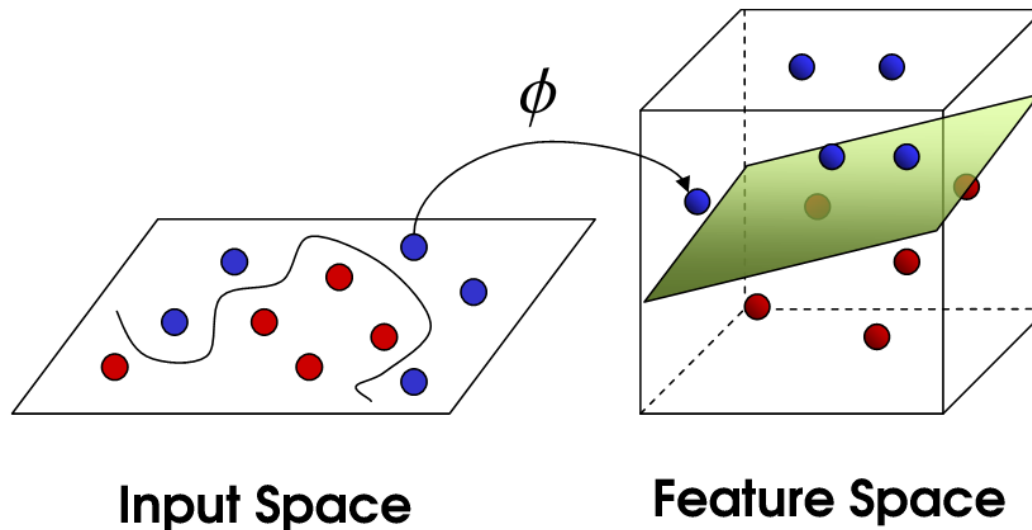


CS109 – Data Science

SVM, Performance evaluation

Joe Blitzstein, Hanspeter Pfister, Verena Kaynig-Fittkau

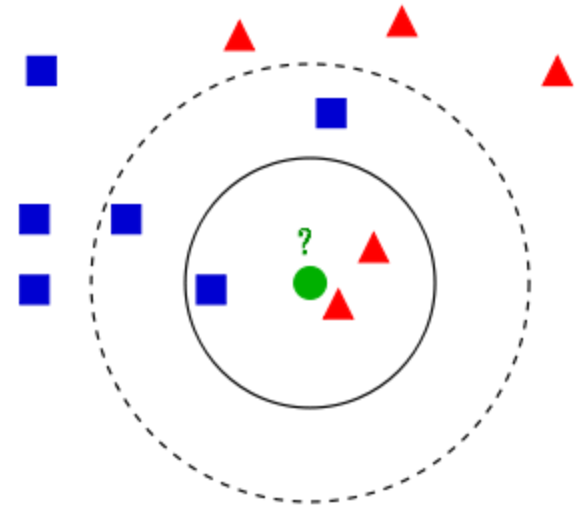


Announcements

- HW1 grades went out yesterday
- They are looking really good, well done everyone!
- HW2 is due this Thursday!
- You should submit an executed notebook
- But please without pages of test output

Recap K-NN

- Keeps all training data
- Training is fast
- Prediction is slow

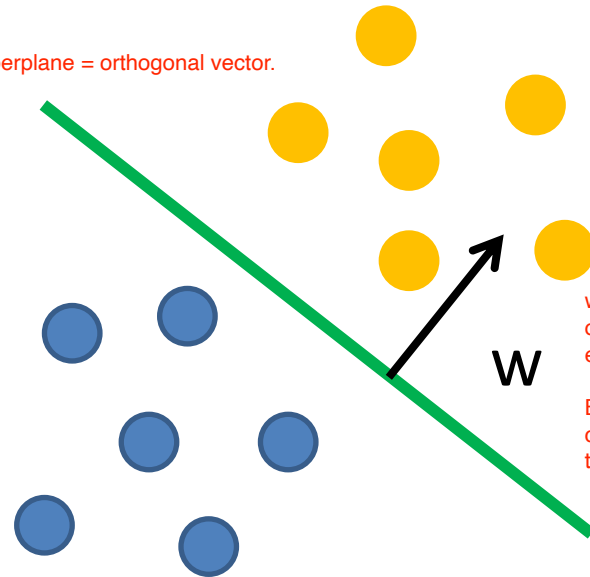


Separating Hyperplane

the goal is to estimate w

- x : data point
- y : label $\in \{-1, +1\}$
- w : weight vector

perpendicular to the hyperplane = orthogonal vector.



w defines the orientation of the hyperplane, in this example.

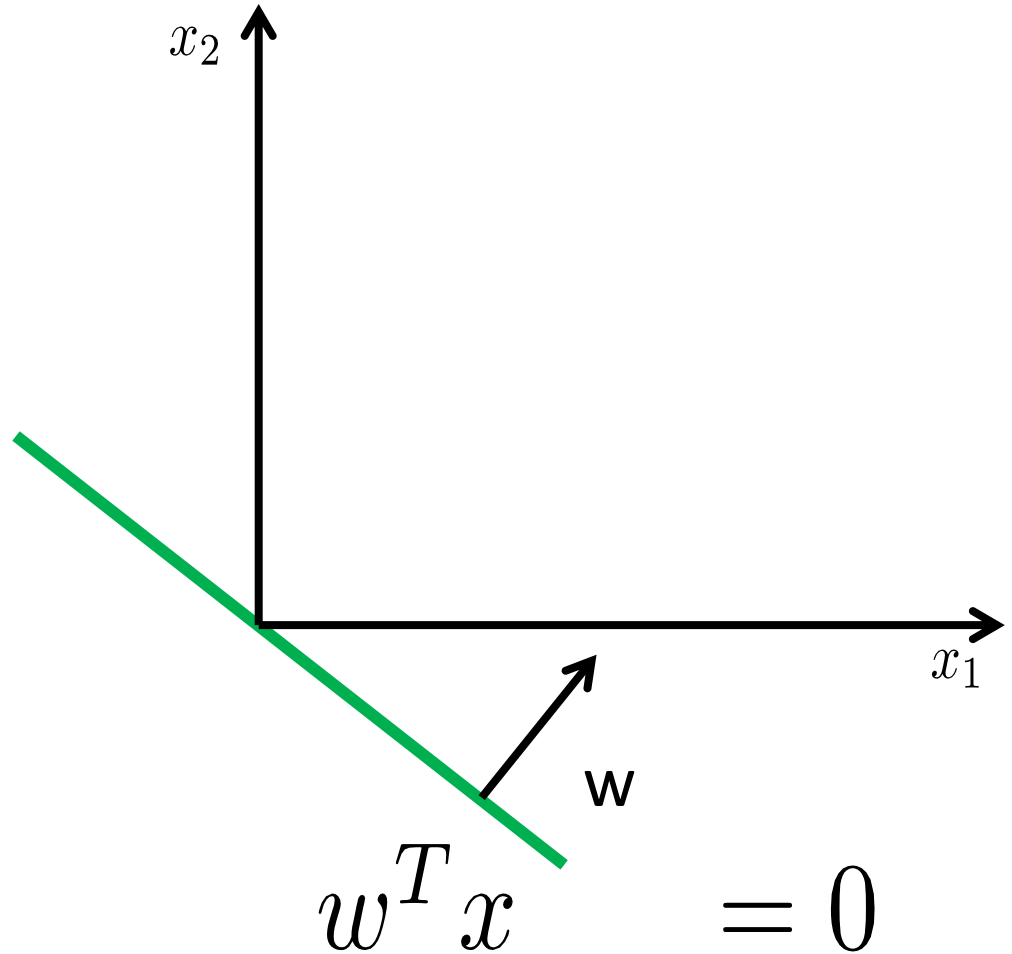
By changing w , you change the orientation of the hyperplane.

Here, the yellow colors are considered to have a label of +1, while blue has the label of -1.

$$w^T x = 0$$

Separating Hyperplane

- x : data point
- y : label $\in \{-1, +1\}$
- w : weight vector




Separating Hyperplane

- x : data point
- y : label $\in \{-1, +1\}$
- w : weight vector
- b : bias

Two things needed to know the separating hyperplane: weight & bias.

weight, to change orientation
bias to change the height of separating hyperplane.

You can then optimize them randomly and have the hyperplane
'wiggle' in between the two classes/labels.



$$w^T x + b = 0$$

Separating Hyperplane


- x : data point
- y : label $\in \{-1, +1\}$
- w : weight vector
- b : bias

Once you have w & b , prediction becomes much easier.

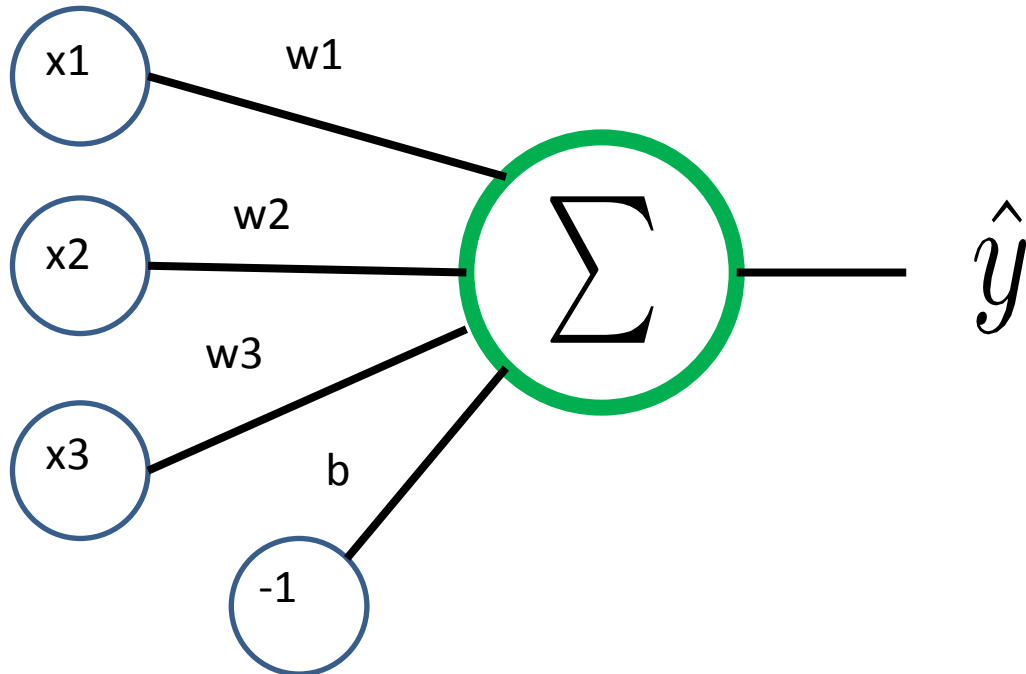
Here, you take your data point x and you evaluate this equation: $w^T x + b$

You look at the sign of that result. If it's greater than 0, then it's on the right side of the hyperplane, and if it's less than 0, it's on the other side of the hyperplane.

All you'd need to store, rather than the training data like in KNN, are the above parameters.




Perceptron



$$w^T x + b = 0$$

Perceptron



Perceptron History

- invented 1957
- by Frank Rosenblatt
- the embryo of an electronic computer that [the Navy] expects will be able to walk, talk, see, write, reproduce itself and be conscious of its existence. (NYT 1958)

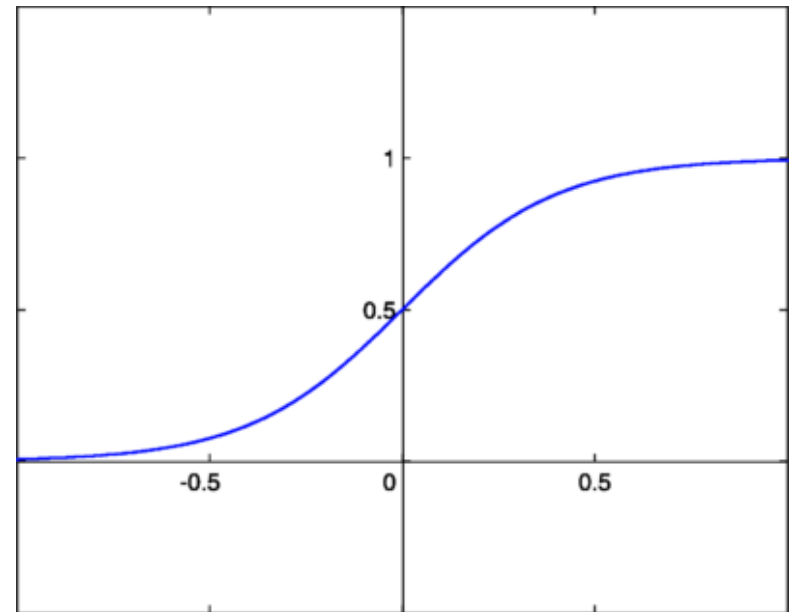
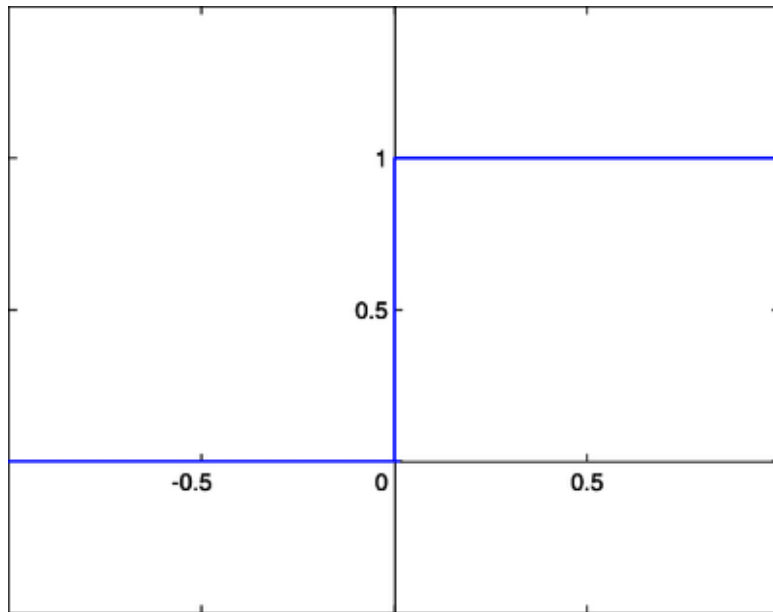
(<http://en.wikipedia.org/wiki/Perceptron>)



Perceptron.mp4

https://www.youtube.com/watch?v=cNxadbrN_al&list=PLdVOMWcqwwllaygvb9ZteZ1r4Br6kRuBO

Side Note: Step vs Sigmoid Activation




$$s(x) = \frac{1}{1 + e^{-cx}}$$

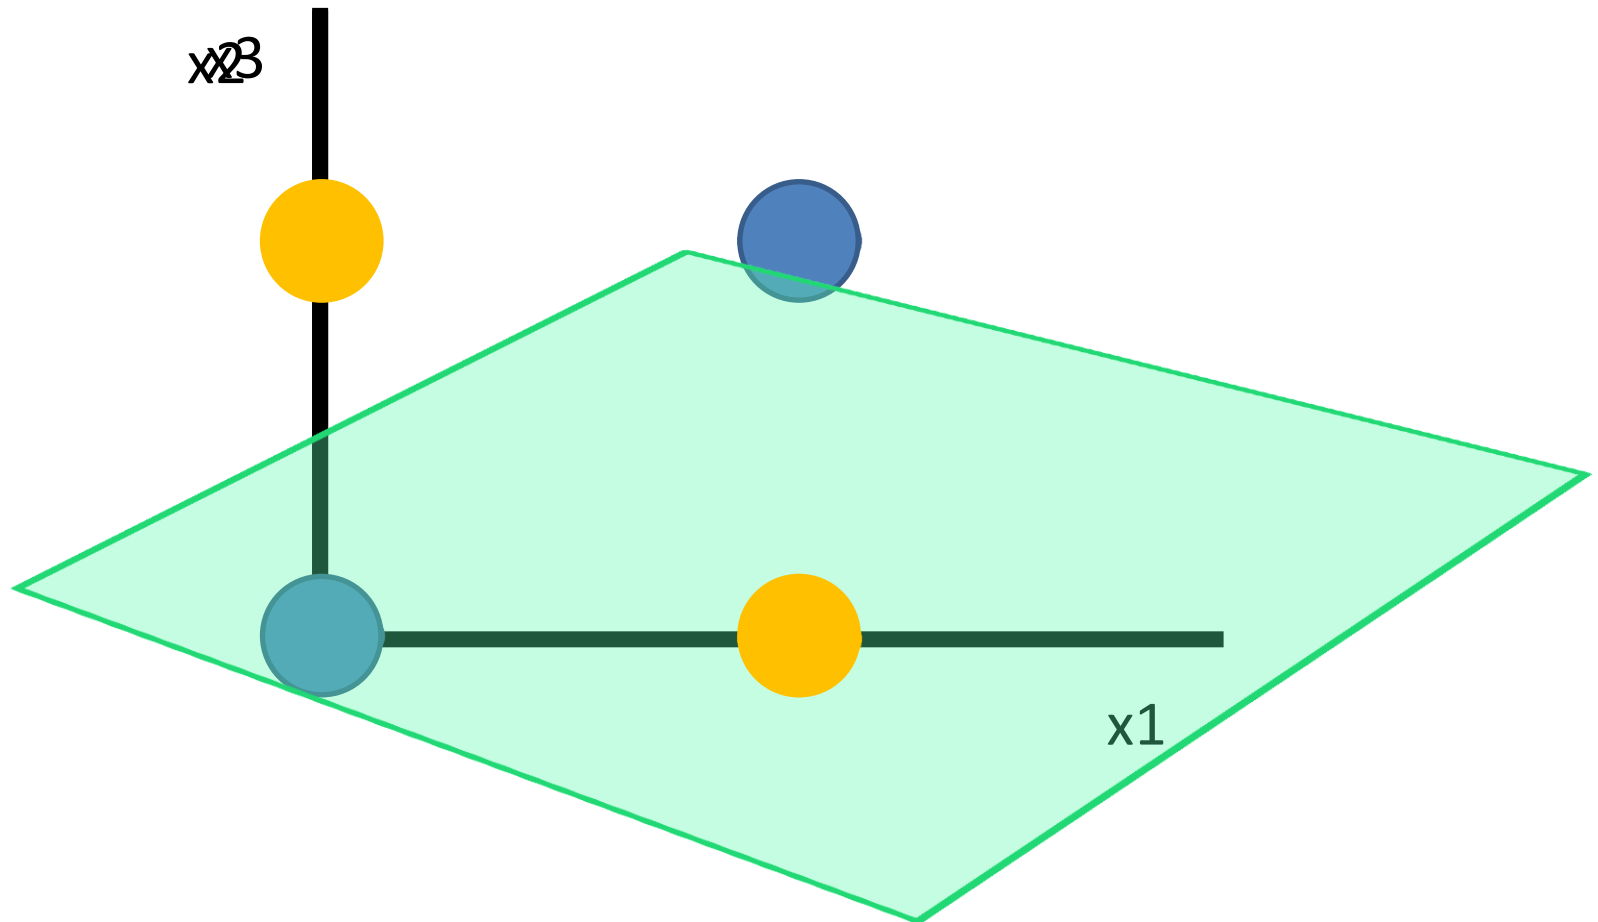
The Critics

- 1969: Minsky and Papert publish their book “Perceptrons”
- Very controversial book, some blame the book for causing the whole research area to stagnate.

The XOR Problem




The XOR Problem



Support Vector Machine

- Widely used for all sorts of classification problems
- Some people say it is the best of the shelf classifier out there

Maximum Margin Classification



Solution depends only on the support vectors!

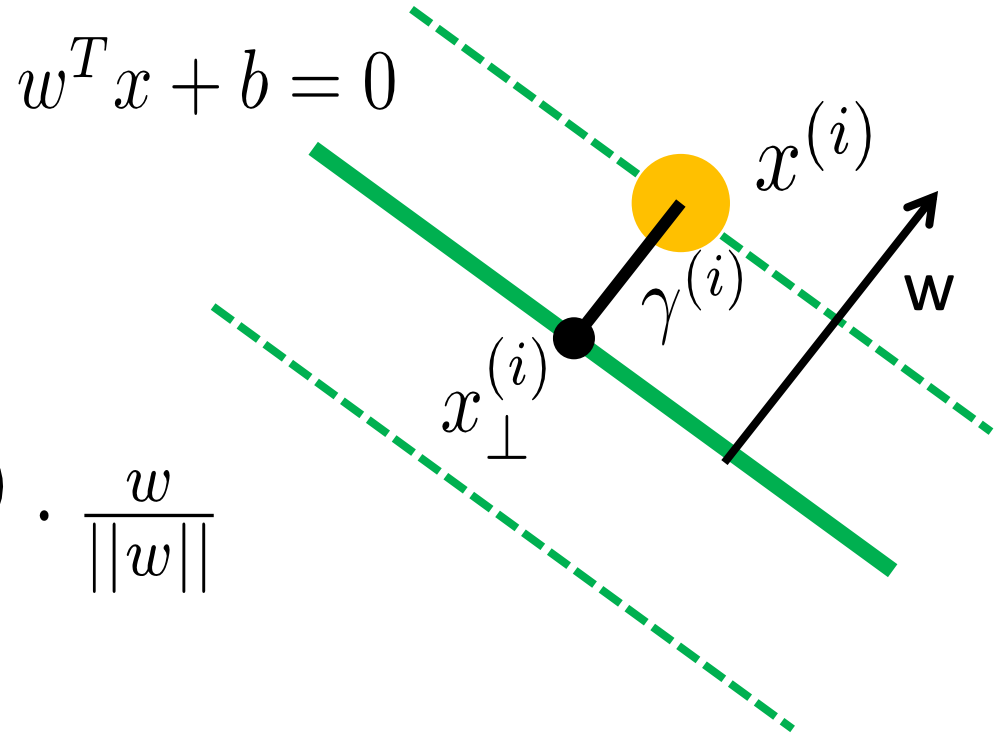
Maximum Margin Classification

margin:

$$x_{\perp}^{(i)} = x^{(i)} - \gamma^{(i)} \cdot \frac{w}{||w||}$$

$$w^T x_{\perp}^{(i)} + b = 0$$

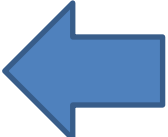
➡ $\gamma^{(i)} = \left(\frac{w^T x^{(i)} + b}{||w||} \right)$



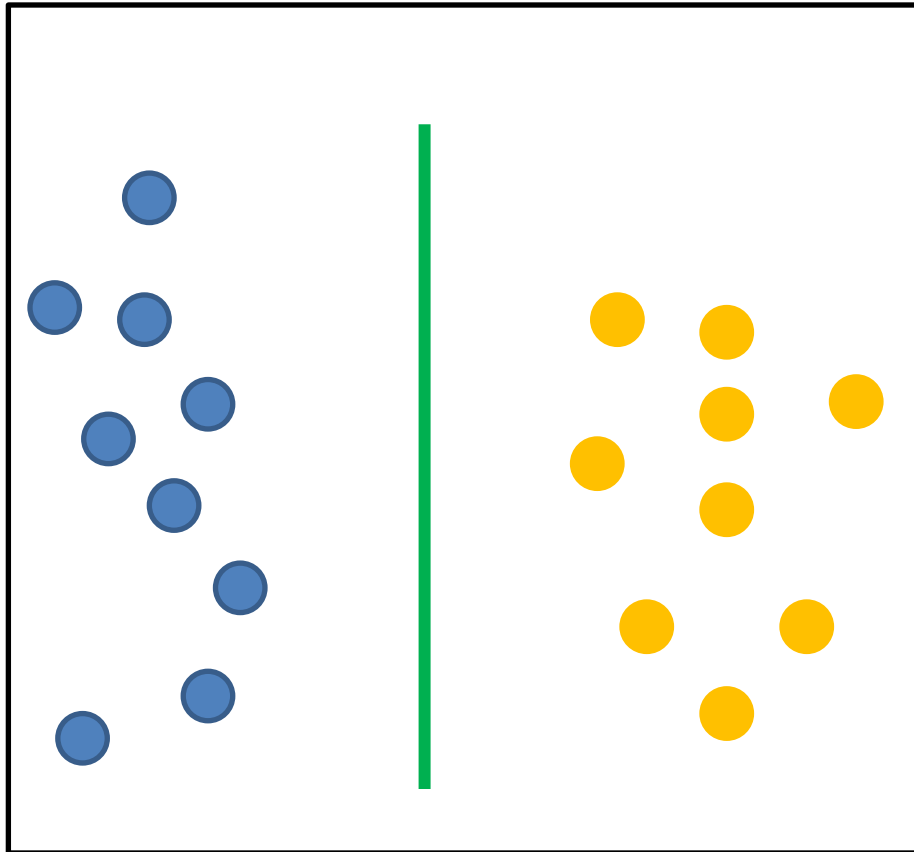
Maximum Margin Classification

$$\gamma^{(i)} = y^{(i)}(w^T x + b)$$

$$\begin{array}{ll} \max_{\gamma, w, b} & \gamma \\ \text{s.t.} & y^{(i)}(w^T x^{(i)} + b) \geq \gamma, \quad i = 1, \dots, m \\ & ||w|| = 1. \end{array}$$

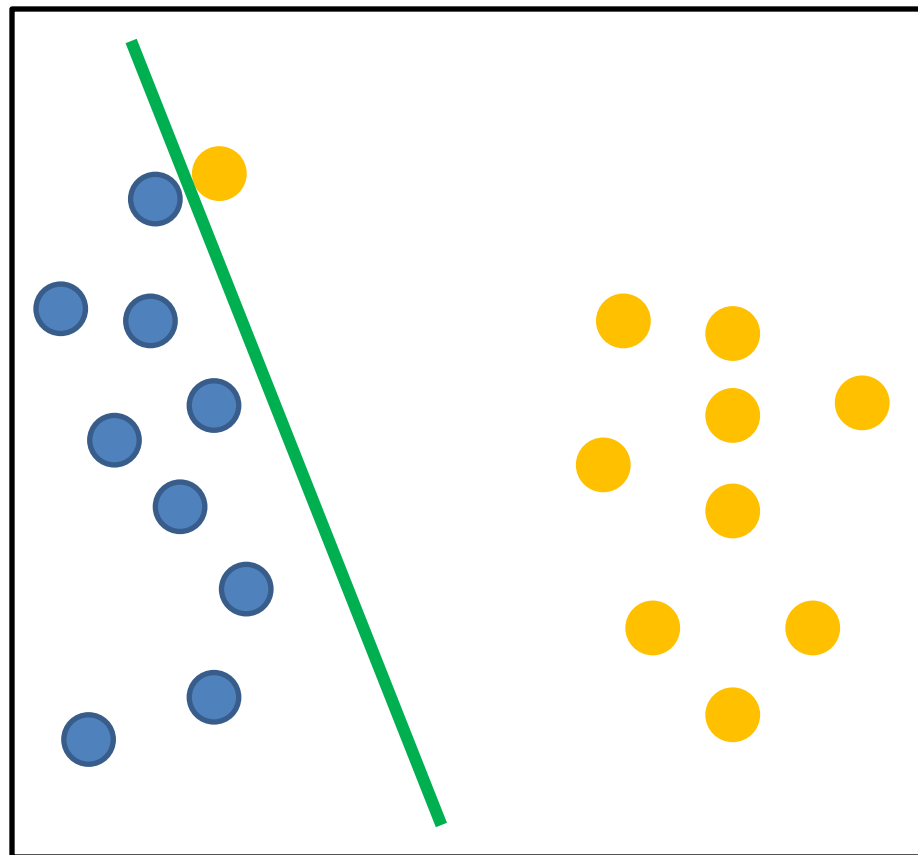
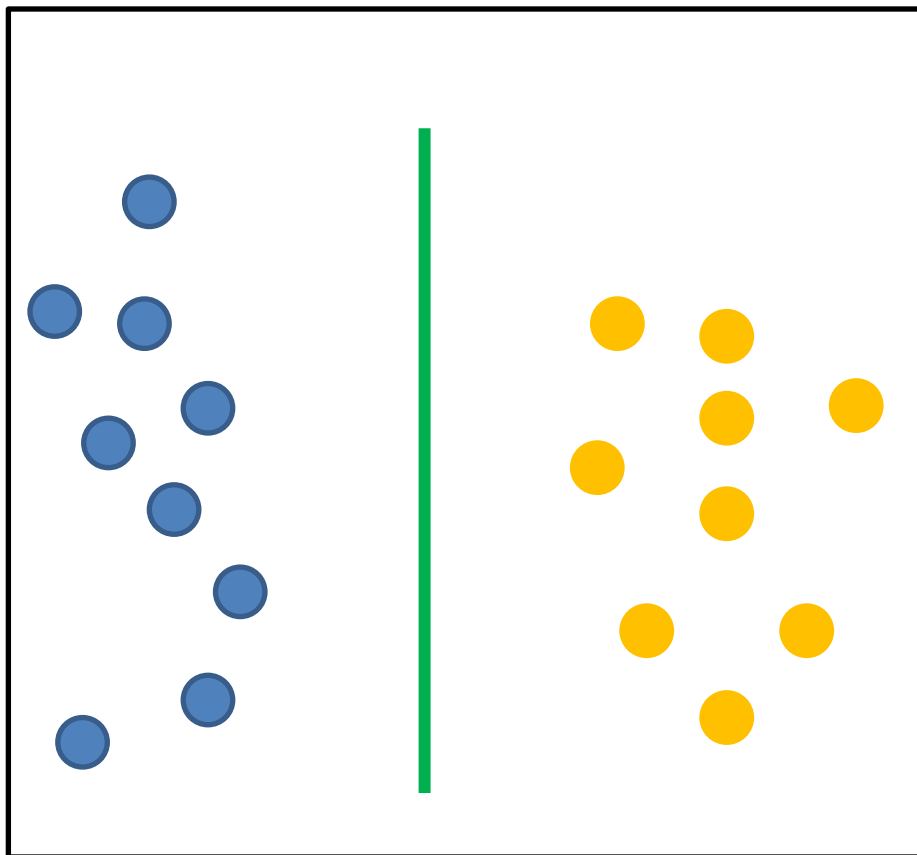
 non-convex

This Is Kind of Odd



- Which data points do we care the most about?
- What would those samples look like?

Two Very Similar Problems



What about outliers?

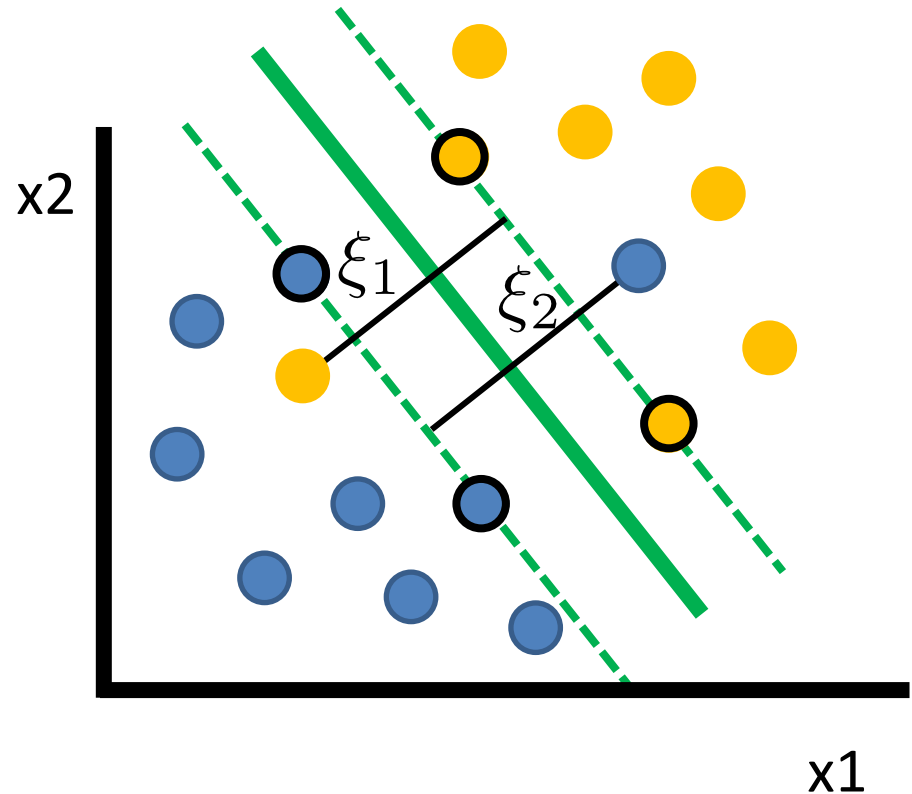
ξ_i : slack variables

$$\min_{w,b,\xi} \frac{1}{2} \|w\|^2$$

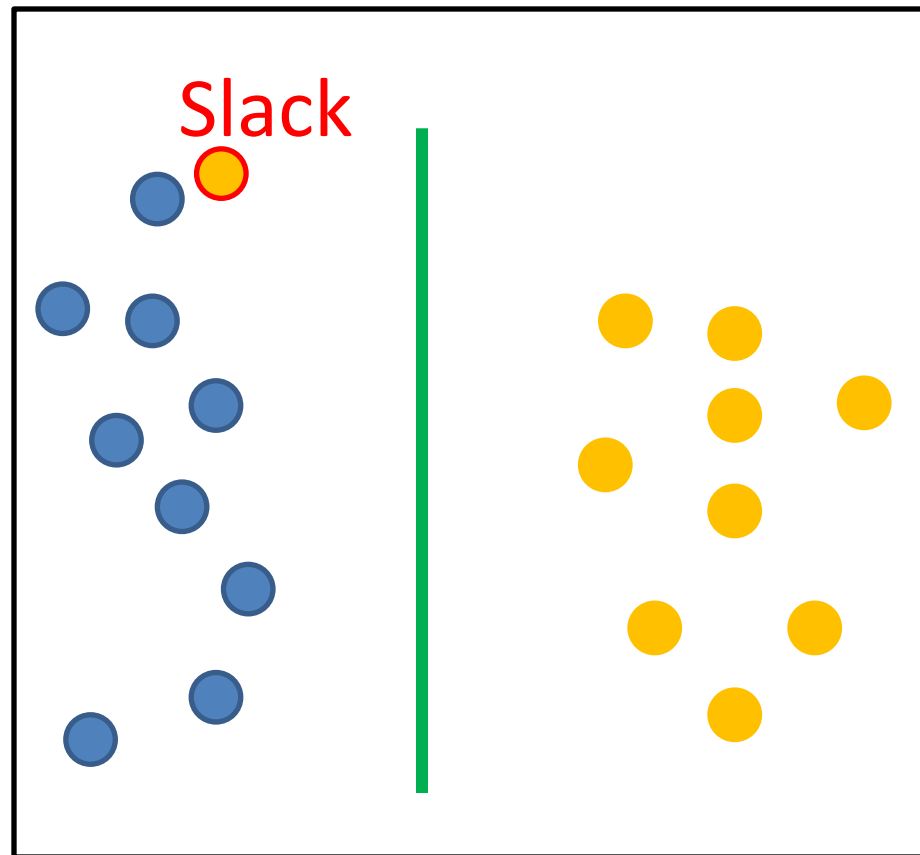
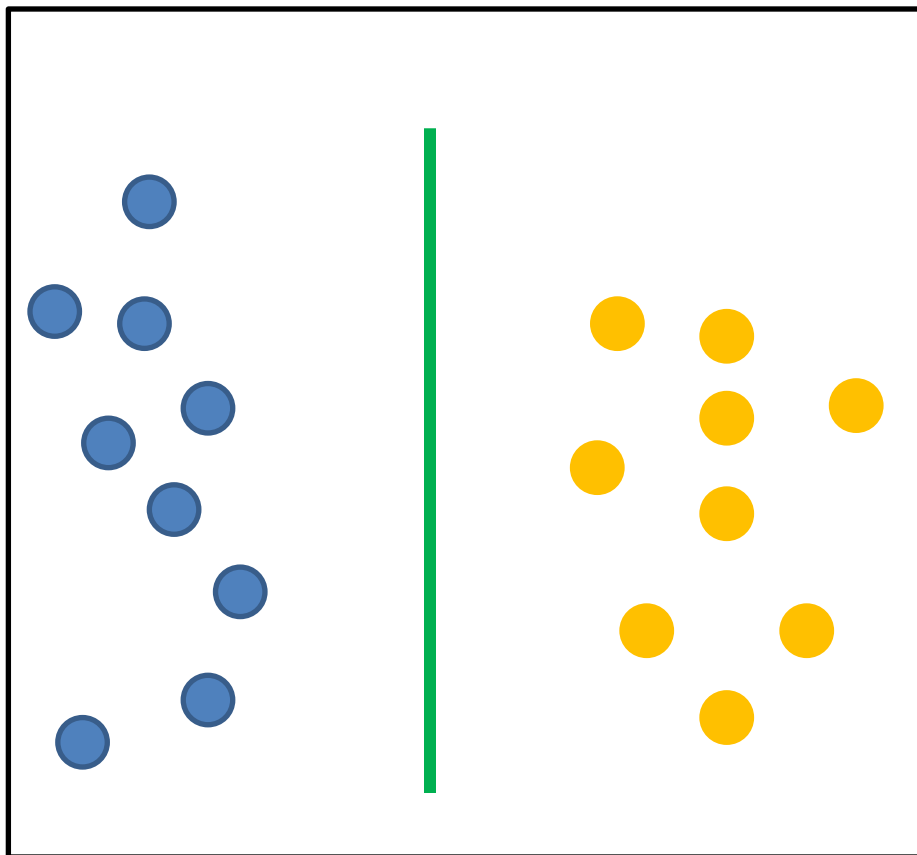
subject to:

$$y^{(i)}(w^T x^{(i)} + b) \geq 1$$

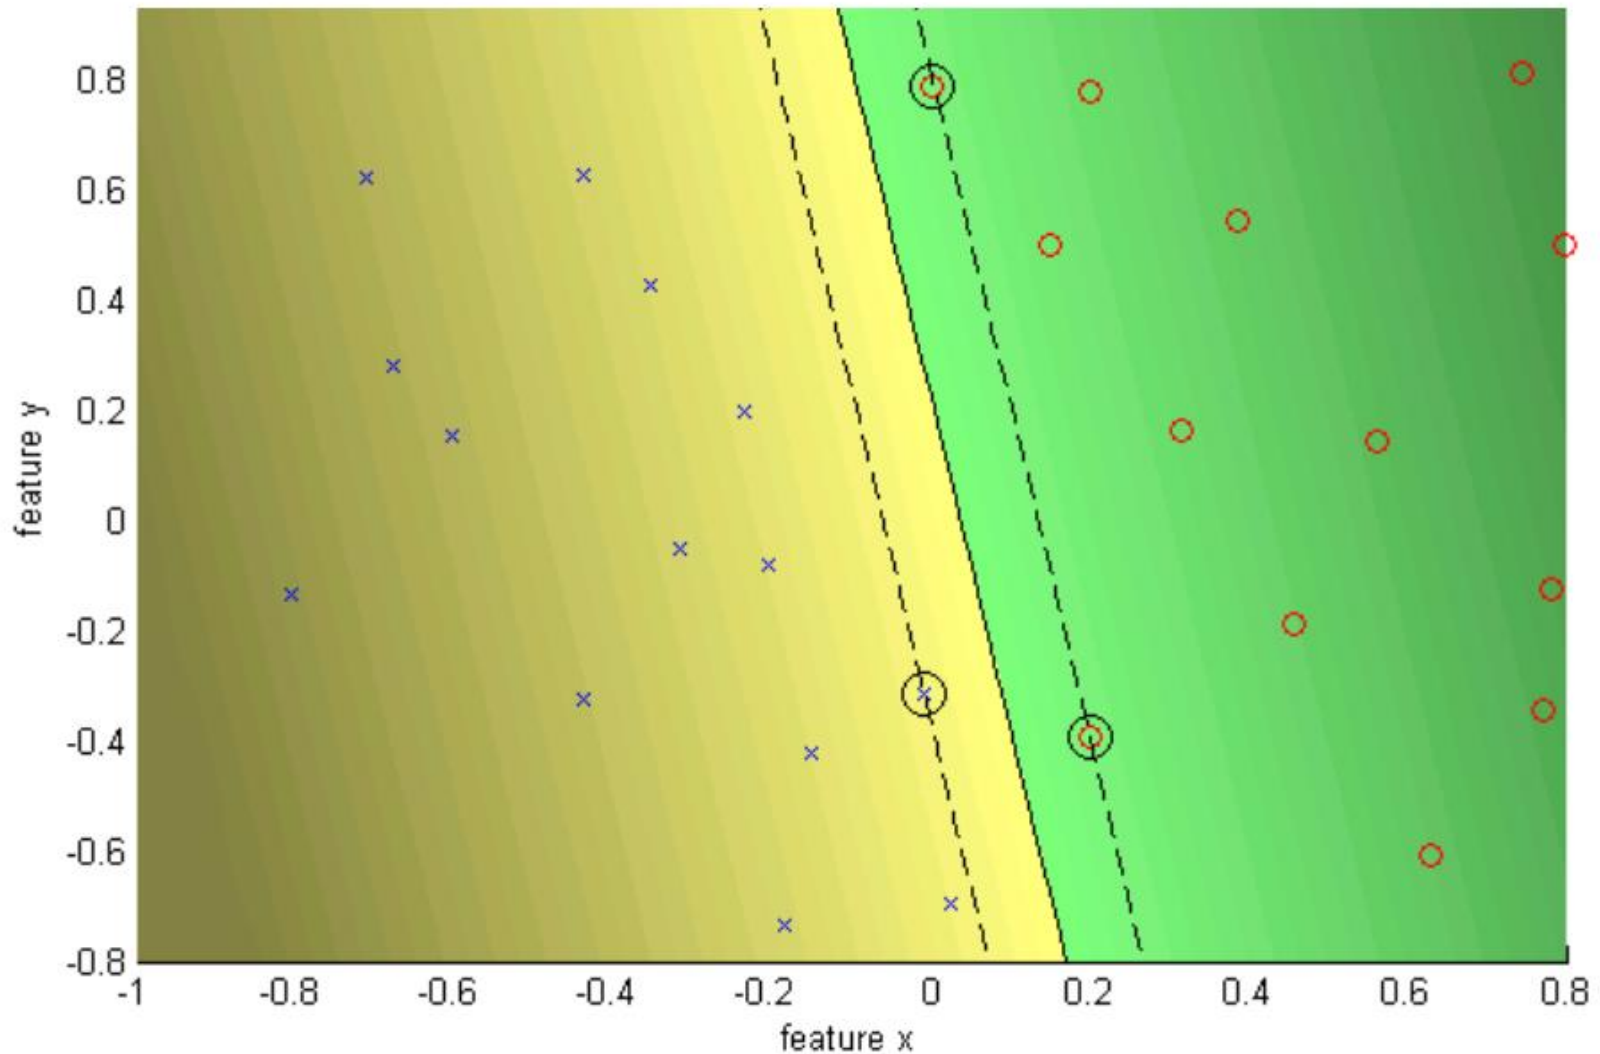
$$(i = 1, \dots, n)$$



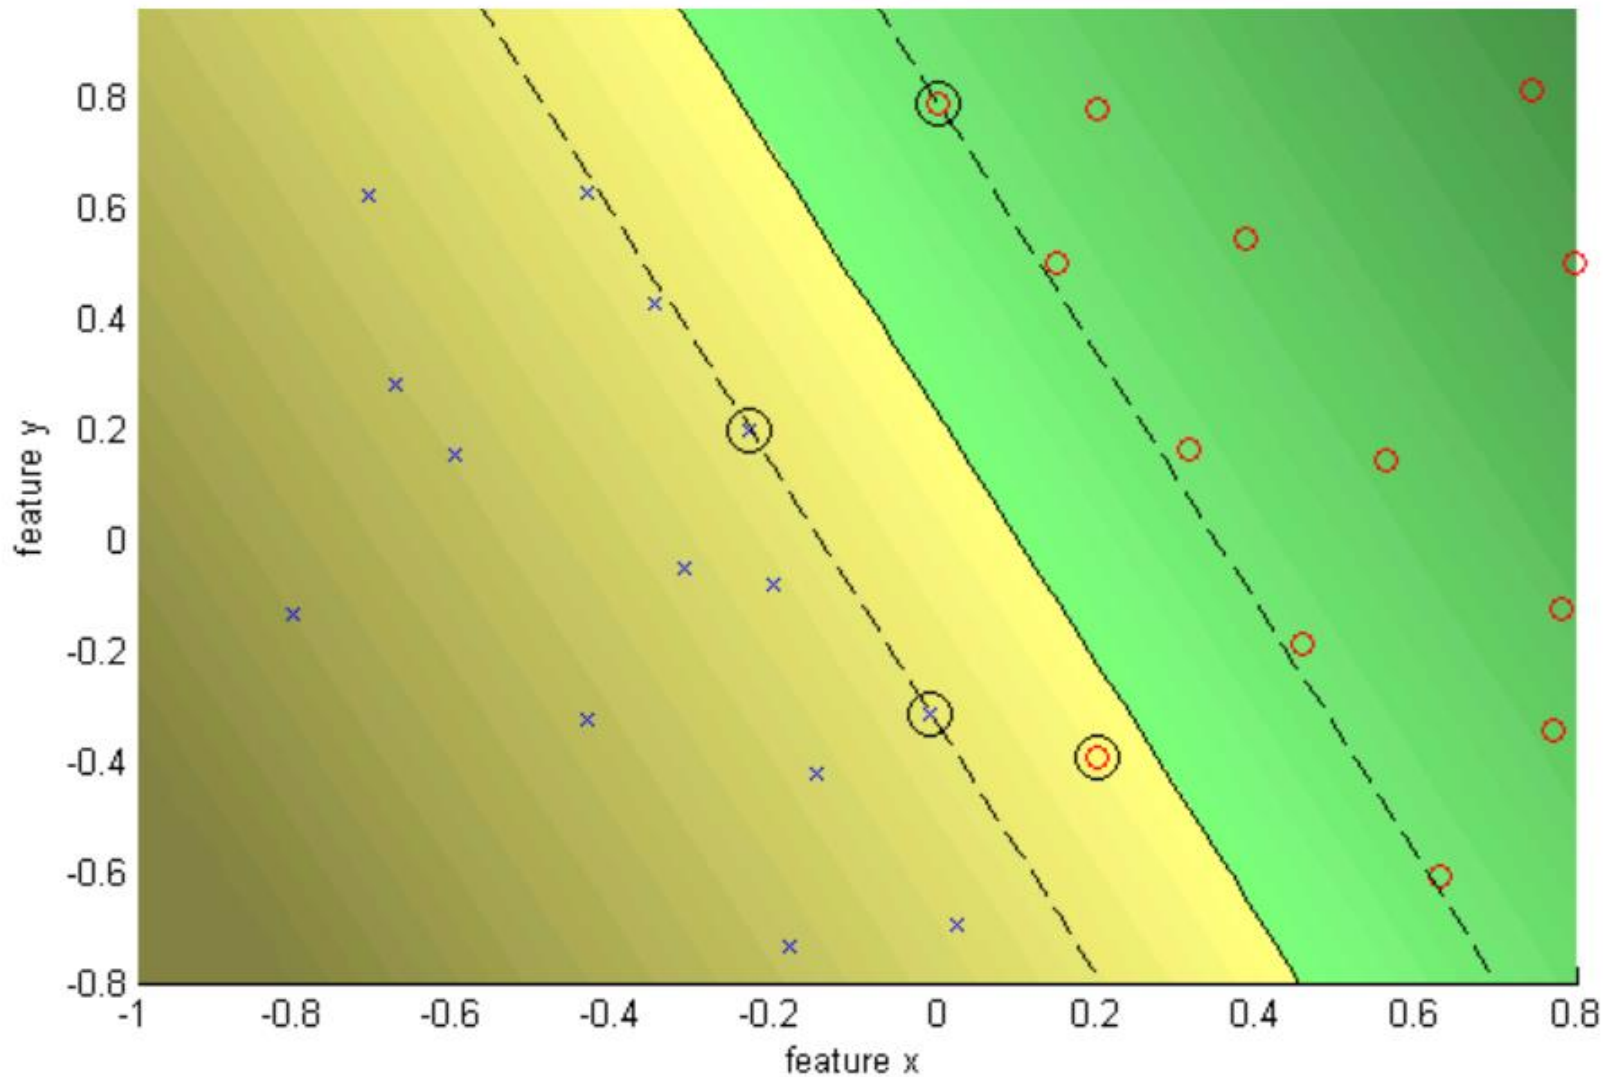
Two Very Similar Problems



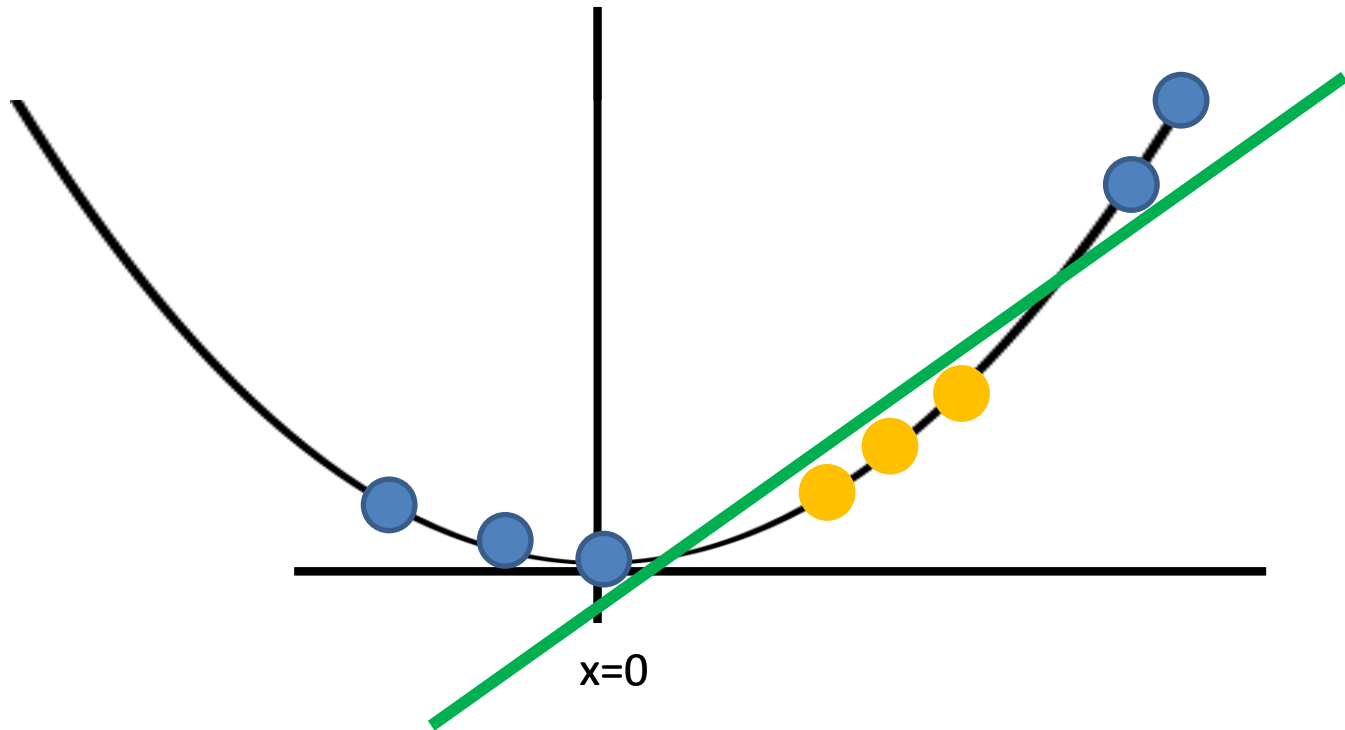
Hard Margin ($C = \text{Infinity}$)



Soft Margin ($C = 10$)

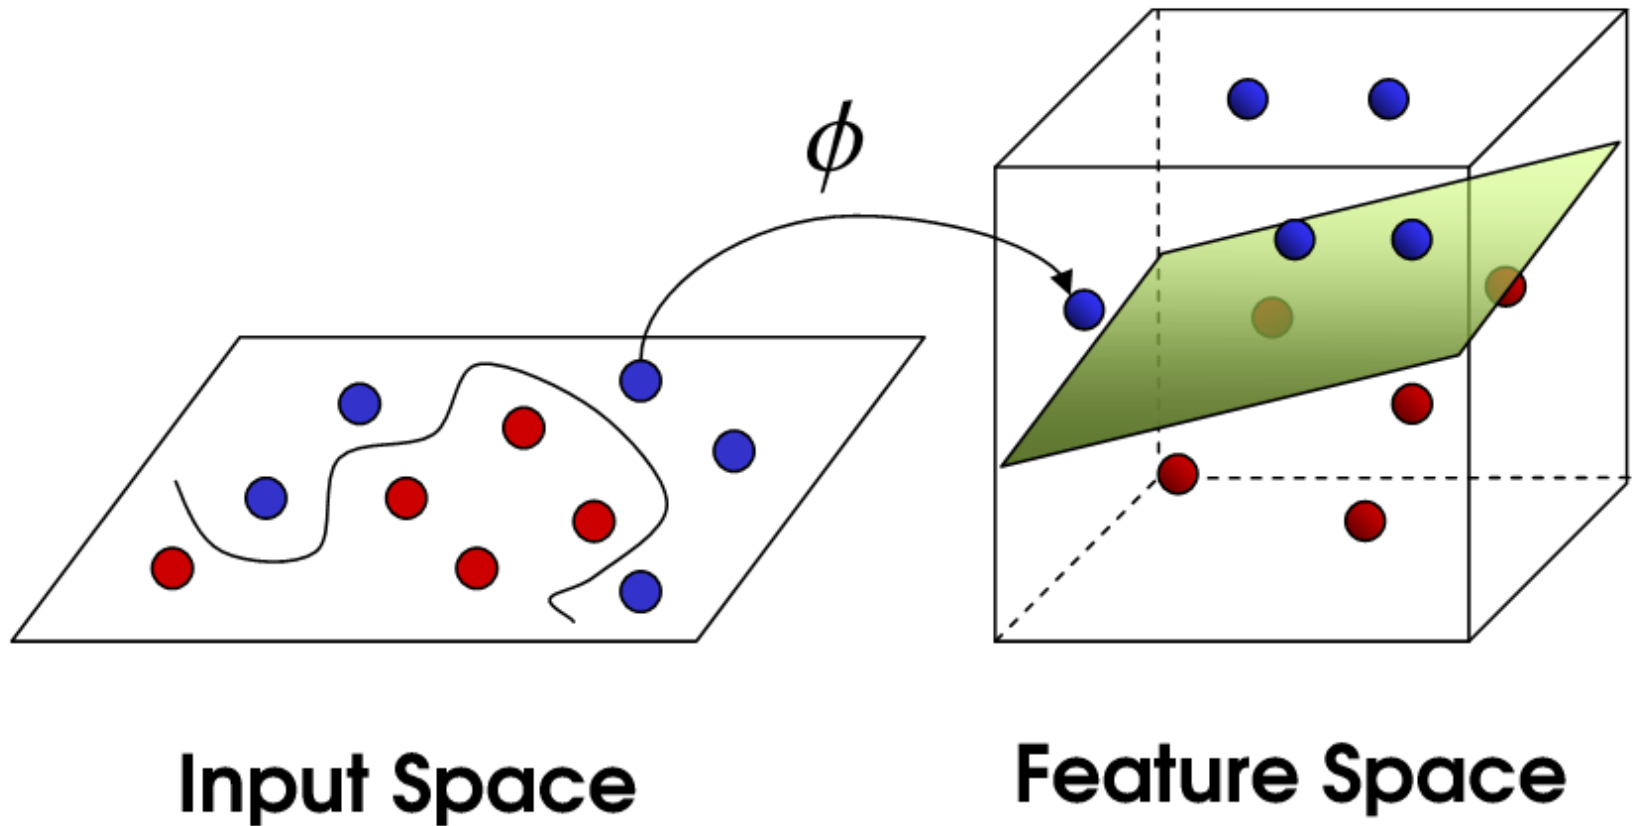


XOR problem revised



Did we add information to make the problem separable?

Non-Linear Decision Boundary



SVM with a polynomial Kernel visualization

Created by:
Udi Aharoni

Quadratic Kernel

$$x = (x_1, x_2)$$

$$\Phi(x) = (1, \sqrt{2}x_1, \sqrt{2}x_2, x_1^2, x_2^2, \sqrt{2}x_1x_2)$$

$$\begin{aligned}\Phi(x) \cdot \Phi(z) &= 1 + 2 \sum_{i=1}^d x_i z_i \\ &\quad + \sum_{i=1}^d x_i^2 z_i^2 + 2 \sum_{i=1}^d \sum_{j=i+1}^d x_i x_j z_i z_j\end{aligned}$$

$$= (1 + x \cdot z)^2$$

Kernel Functions

$$K(x, z) = \Phi(x) \cdot \Phi(z)$$

- Polynomial:

$$K(x, z) = (1 + x \cdot z)^s$$

- Radial basis function (RBF):

$$K(x, z) = \exp(-\gamma(x - z)^2)$$

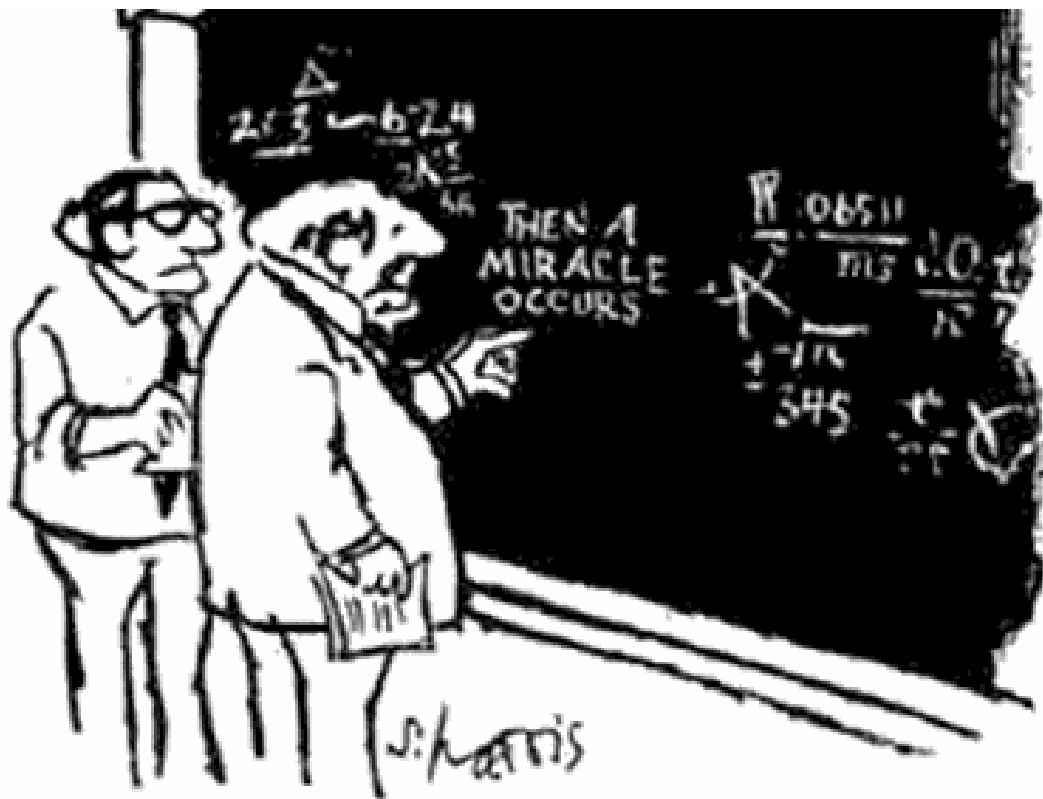
So what is the excitement?

$$\max_{\alpha} \sum$$

$$\text{s.t. } \alpha_i$$

$$\sum$$

$$(i)^T x(j)$$

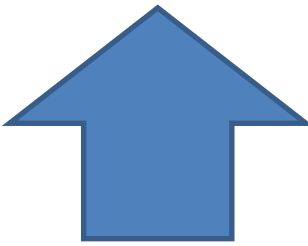


$$\arg r$$


s.t. y "I THINK YOU SHOULD BE MORE EXPLICIT
HERE IN STEP TWO."

So what is the excitement?

$$\begin{aligned}
 & \max_{\alpha} \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i,j=1}^m y^{(i)} y^{(j)} \alpha_i \alpha_j \boxed{x^{(i)T} x^{(j)}} \\
 & \text{s.t. } \alpha_i \geq 0, \quad i = 1, \dots, m \\
 & \quad \sum_{i=1}^m \alpha_i y^{(i)} = 0
 \end{aligned}$$



$$\begin{aligned}
 & \arg \min_{w,b} \frac{1}{2} ||w||^2 \\
 & \text{s.t. } y^{(i)} (w^T x^{(i)} + b) \geq 1
 \end{aligned}$$



 $\boxed{K(x^{(i)}, x^{(j)})}$

Prediction

$$w^T x + b = \sum_{i=1}^m \alpha_i y^{(i)} \langle x^{(i)}, x \rangle + b.$$

- Again we can use the kernel trick!
- Prediction speed depends on number of support vectors

The Miracle Explained

- Andrew Ng does this really well
- <http://cs229.stanford.edu/notes/cs229-notes3.pdf>
- Course is also on Youtube, ItunesU, etc.

Kernel Trick for SVMs

- Arbitrary many dimensions
- Little computational cost
- Maximal margin helps with curse of dimensionality

Face Recognition

pred: Colin_Powell
true: Colin_Powell



pred: George_W_Bush
true: George_W_Bush



pred: Colin_Powell
true: Colin_Powell



pred: Tony_Blair
true: Tony_Blair



pred: George_W_Bush
true: George_W_Bush



pred: Colin_Powell
true: Colin_Powell



pred: George_W_Bush
true: George_W_Bush



pred: George_W_Bush
true: George_W_Bush



pred: Tony_Blair
true: Tony_Blair



pred: Colin_Powell
true: Colin_Powell



pred: George_W_Bush
true: George_W_Bush



pred: Donald_Rumsfeld
true: Donald_Rumsfeld



Face Recognition

- Load image data
 - Put your test data aside
 - Extract Eigenfaces
 - Train SVM
 - Evaluate performance
-
- Red are cross validation steps

http://scikit-learn.org/stable/auto_examples/applications/face_recognition.html#example-applications-face-recognition-py

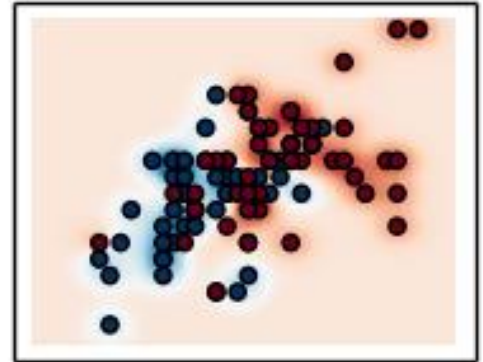
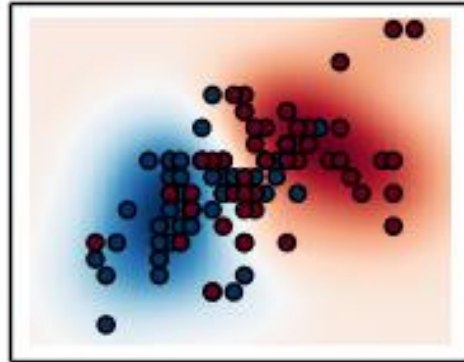
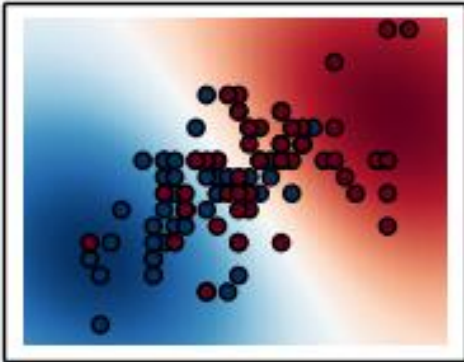


SVM_sign_language.mp4

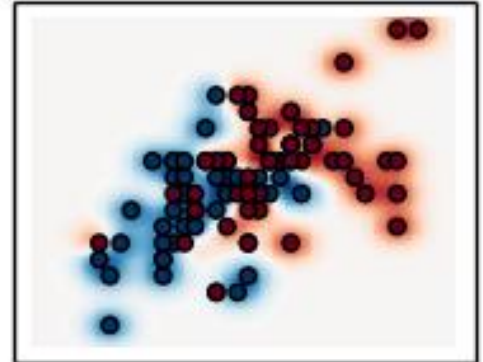
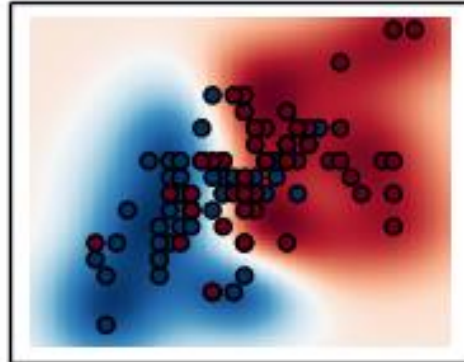
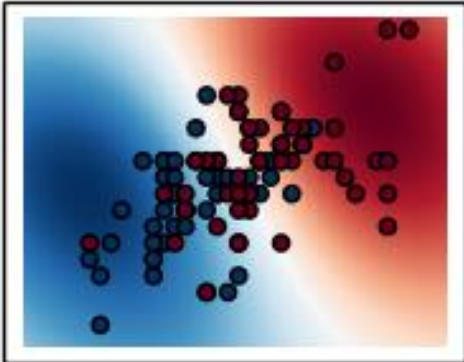
[Jhon Gonzalez](#)

https://www.youtube.com/watch?v=cxHMgl2_5zg

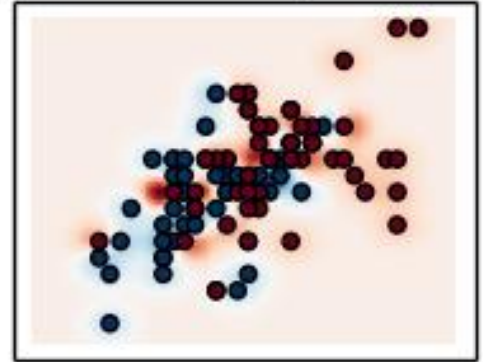
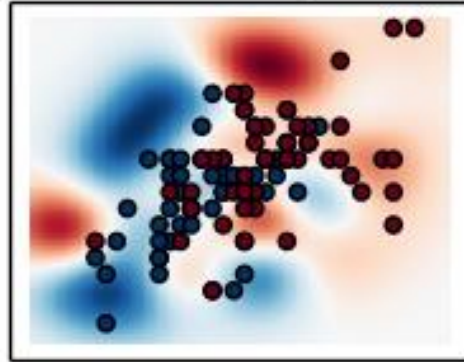
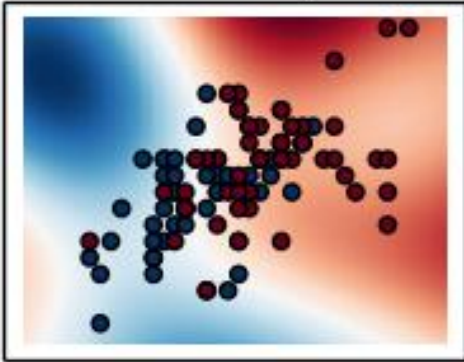
$\gamma=10^{-1}, C=10^{-2}$ $\gamma=10^0, C=10^{-2}$ $\gamma=10^1, C=10^{-2}$



$\gamma=10^{-1}, C=10^0$ $\gamma=10^0, C=10^0$ $\gamma=10^1, C=10^0$



$\gamma=10^{-1}, C=10^2$ $\gamma=10^0, C=10^2$ $\gamma=10^1, C=10^2$



Tips and Tricks

- SVMs are not scale invariant
- Check if your library normalizes by default
- Normalize your data
 - mean: 0 , std: 1
 - map to $[0,1]$ or $[-1,1]$
- Normalize test set in same way!

Tips and Tricks

- RBF kernel is a good default
- For parameters try exponential sequences
- Read:

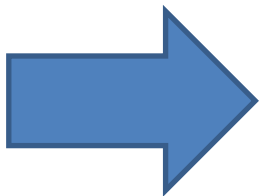
Chih-Wei Hsu et al., “**A Practical Guide to Support Vector Classification**”,
Bioinformatics (2010)

SVM vs KNN

- What are the main key differences?

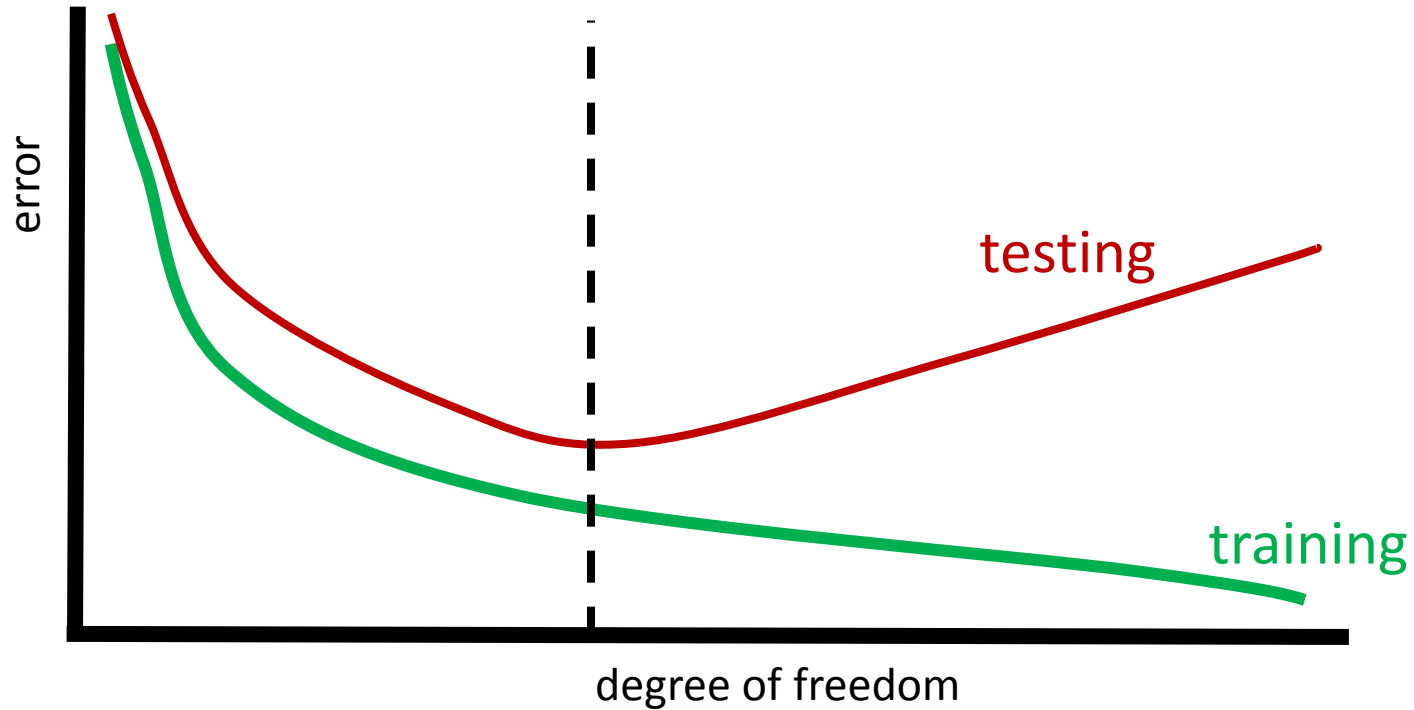
Parameter Tuning

- Given a classification task
- Which kernel ?
- Which kernel parameter values?
- Which value for C?



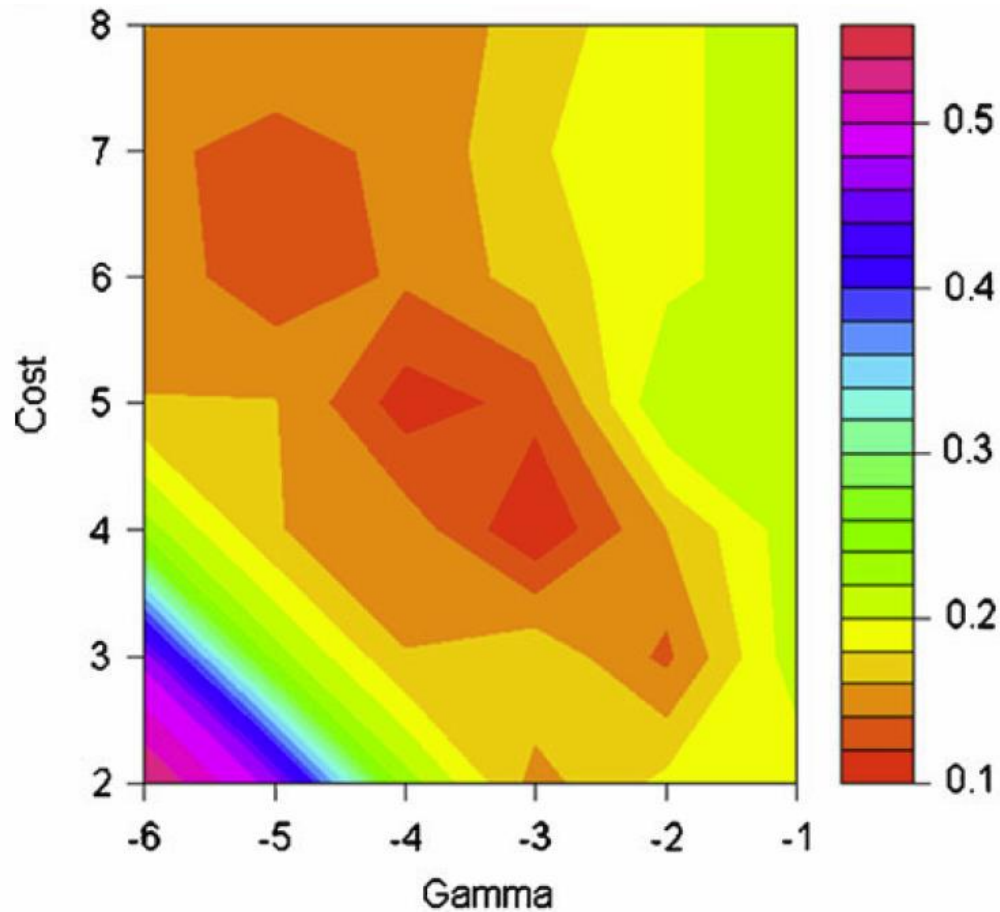
Try different combinations
and take the **best**.

Train vs. Test Error



Where is KNN on this graph for $K=1$, or for $K=\text{Inf}$?

Grid Search



Zang et al., "Identification of heparin samples that contain impurities or contaminants by chemometric pattern recognition analysis of proton NMR spectral data", Anal Bioanal Chem (2011)

Error Measures

- True positive (tp)
- True negative (tn)
- False positive (fp)
- False negative (fn)

		predicted	
		1	-1
true	1	tp	fn
	-1	fp	tn

TPR and FPR

- True Positive Rate:

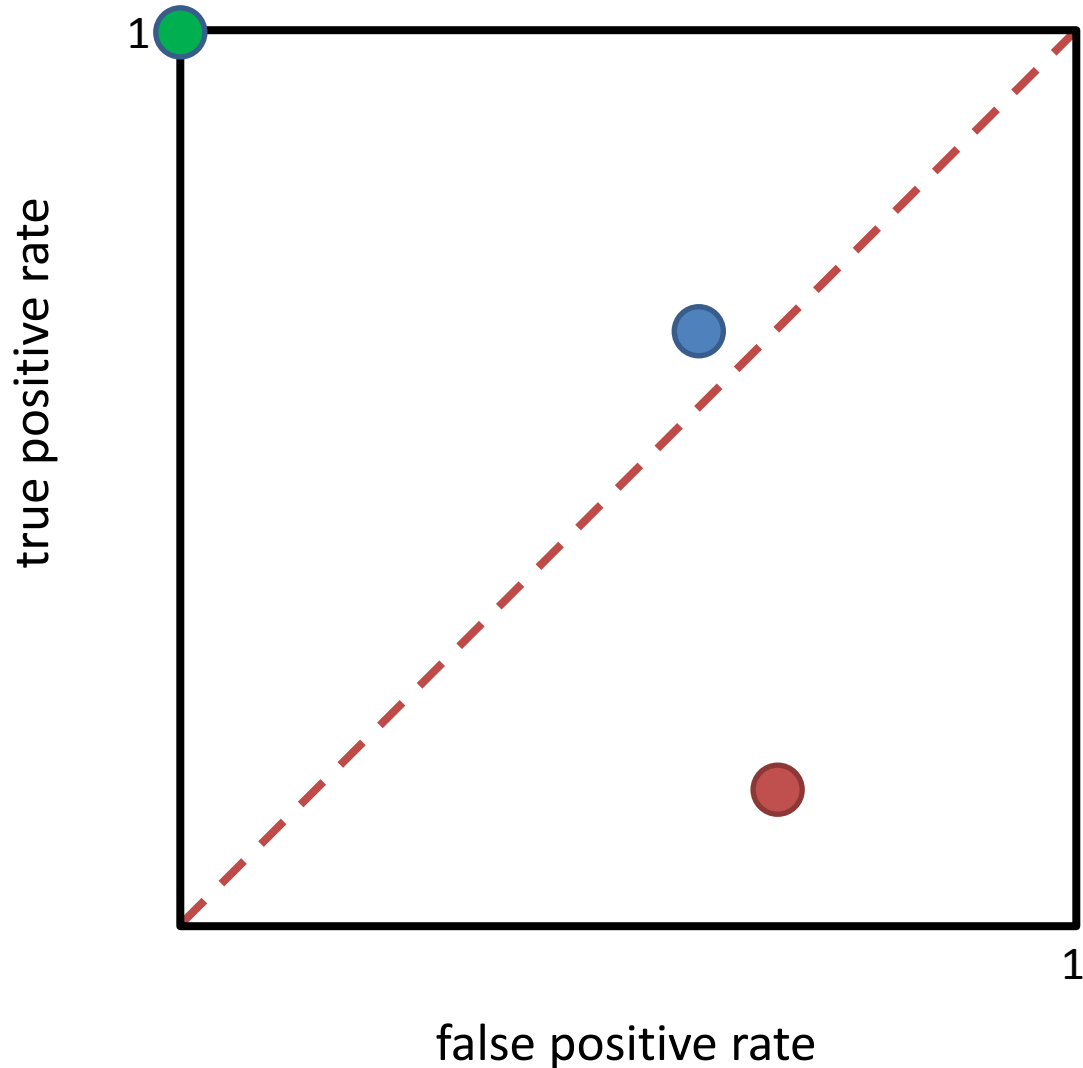
$$\frac{tp}{tp + fn}$$

- False Positive Rate:

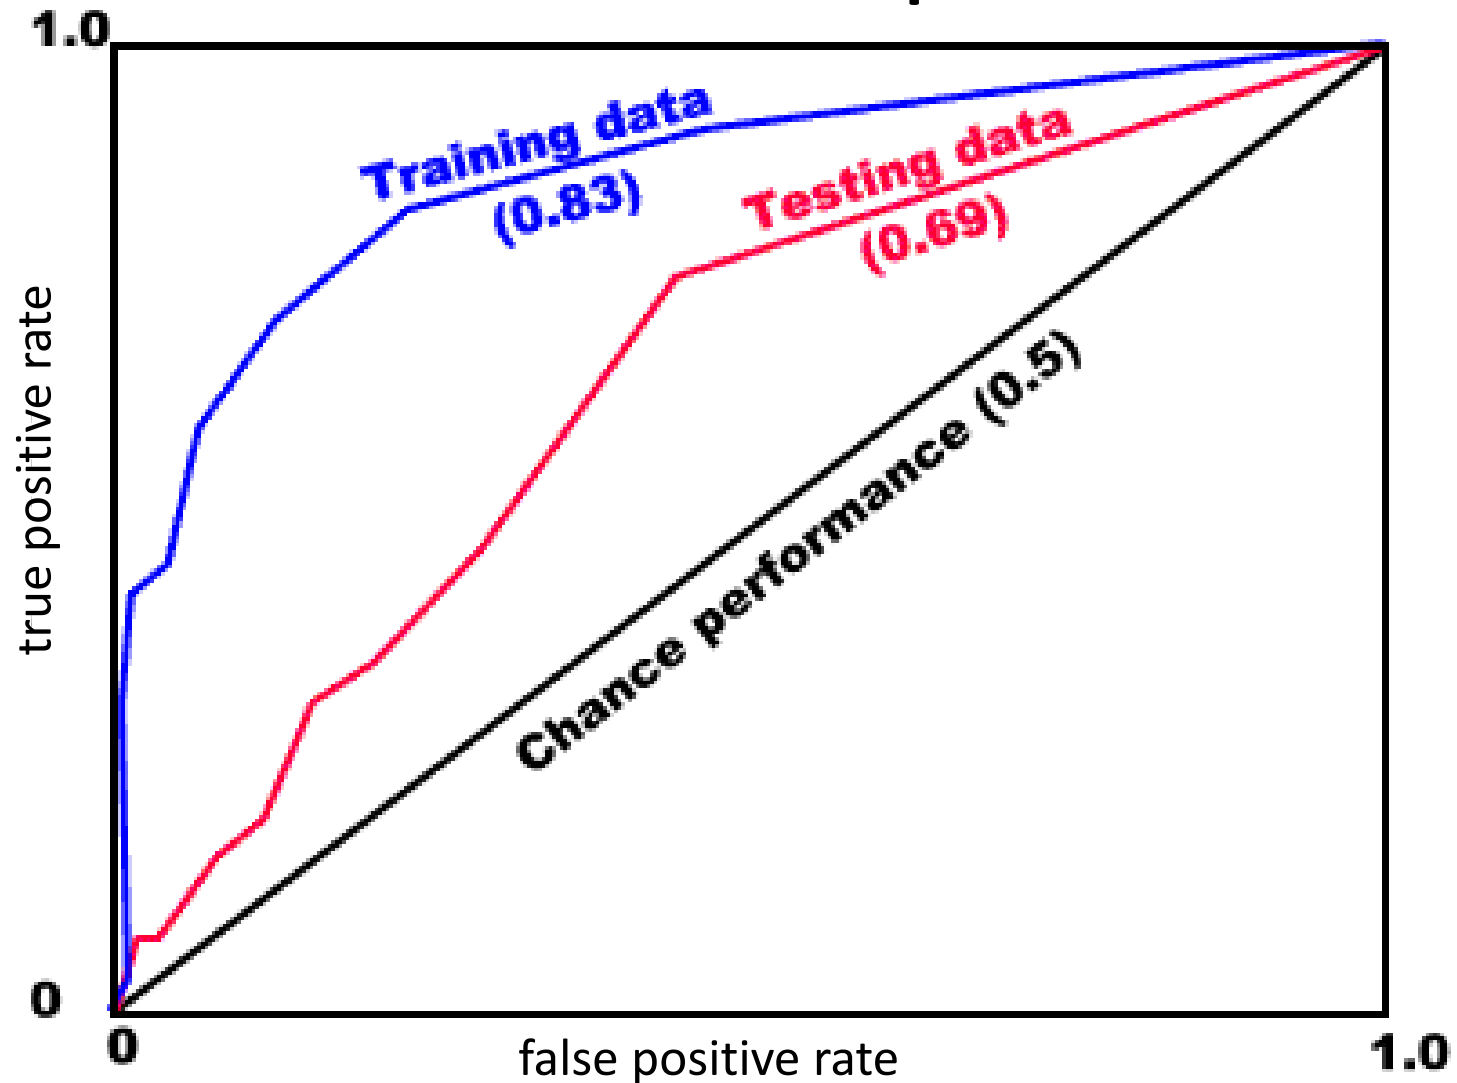
$$\frac{fp}{fp + tn}$$

		predicted	
		1	-1
true	1	tp	fn
	-1	fp	tn

Receiver Operating Characteristic



ROC Example



Precision Recall

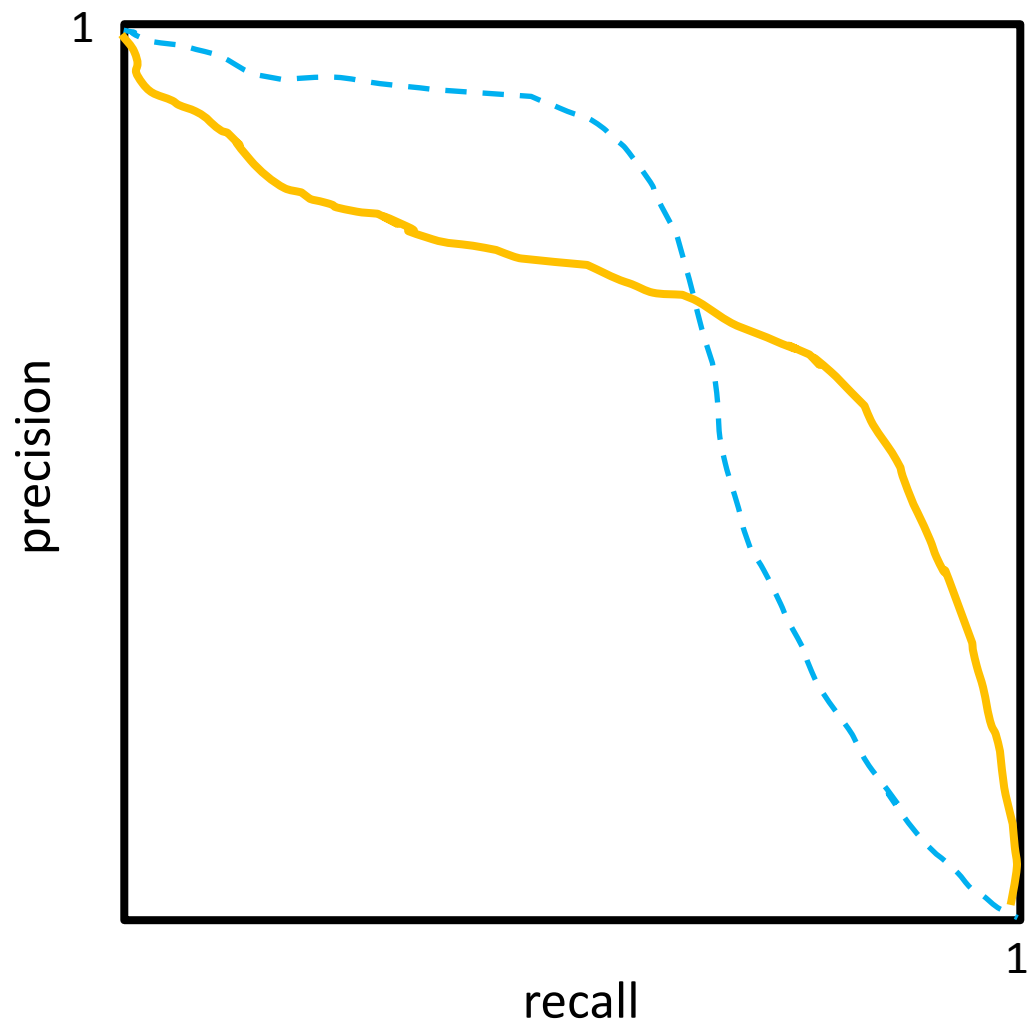
- Recall: $\frac{tp}{tp + fn}$
- Precision: $\frac{tp}{tp + fp}$

		predicted	
		1	-1
true	1	tp	fn
	-1	fp	tn

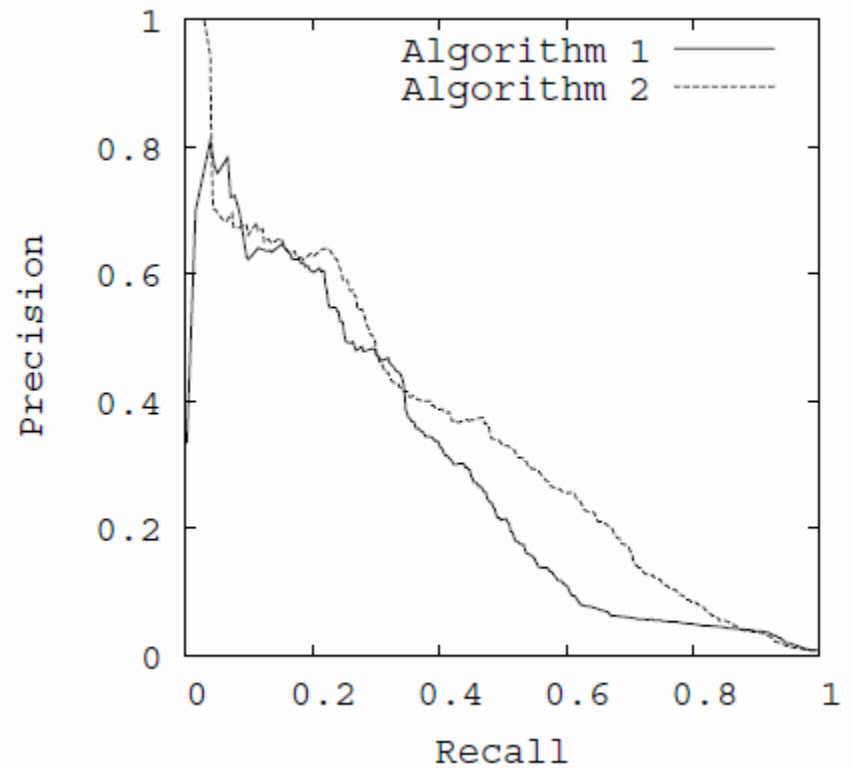
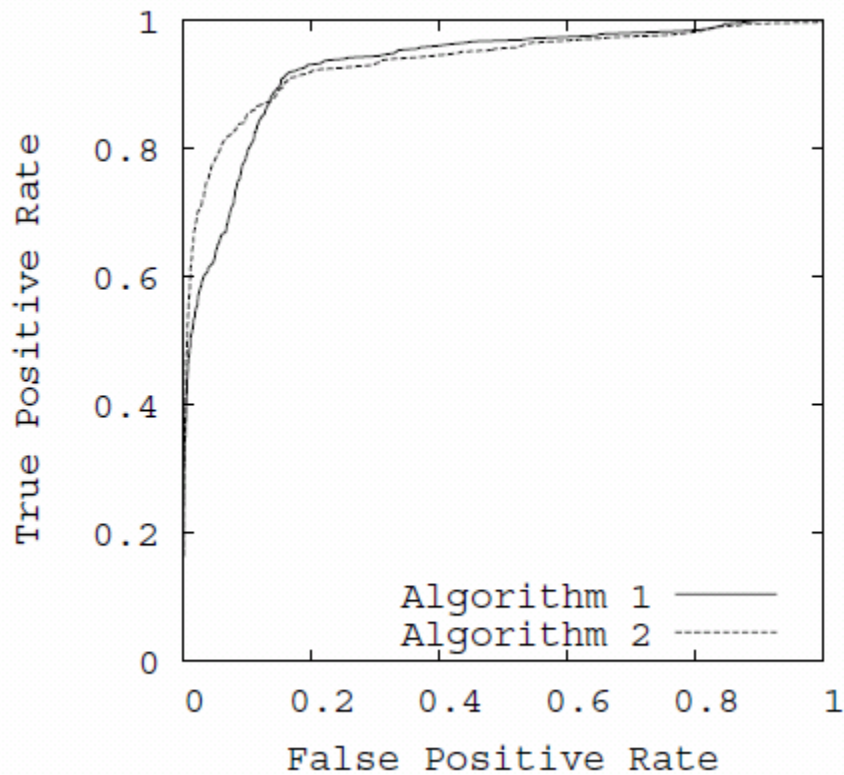
Precision Recall

- **Recall:** If I pick a random positive example, what is the probability of making the right prediction?
- **Precision:** If I take a positive prediction example, what is the probability that it is indeed a positive example?

Precision Recall Curve



Comparison



J. Davis & M. Goadrich,
“The Relationship Between Precision-Recall and ROC Curves.”,
ICML (2006)

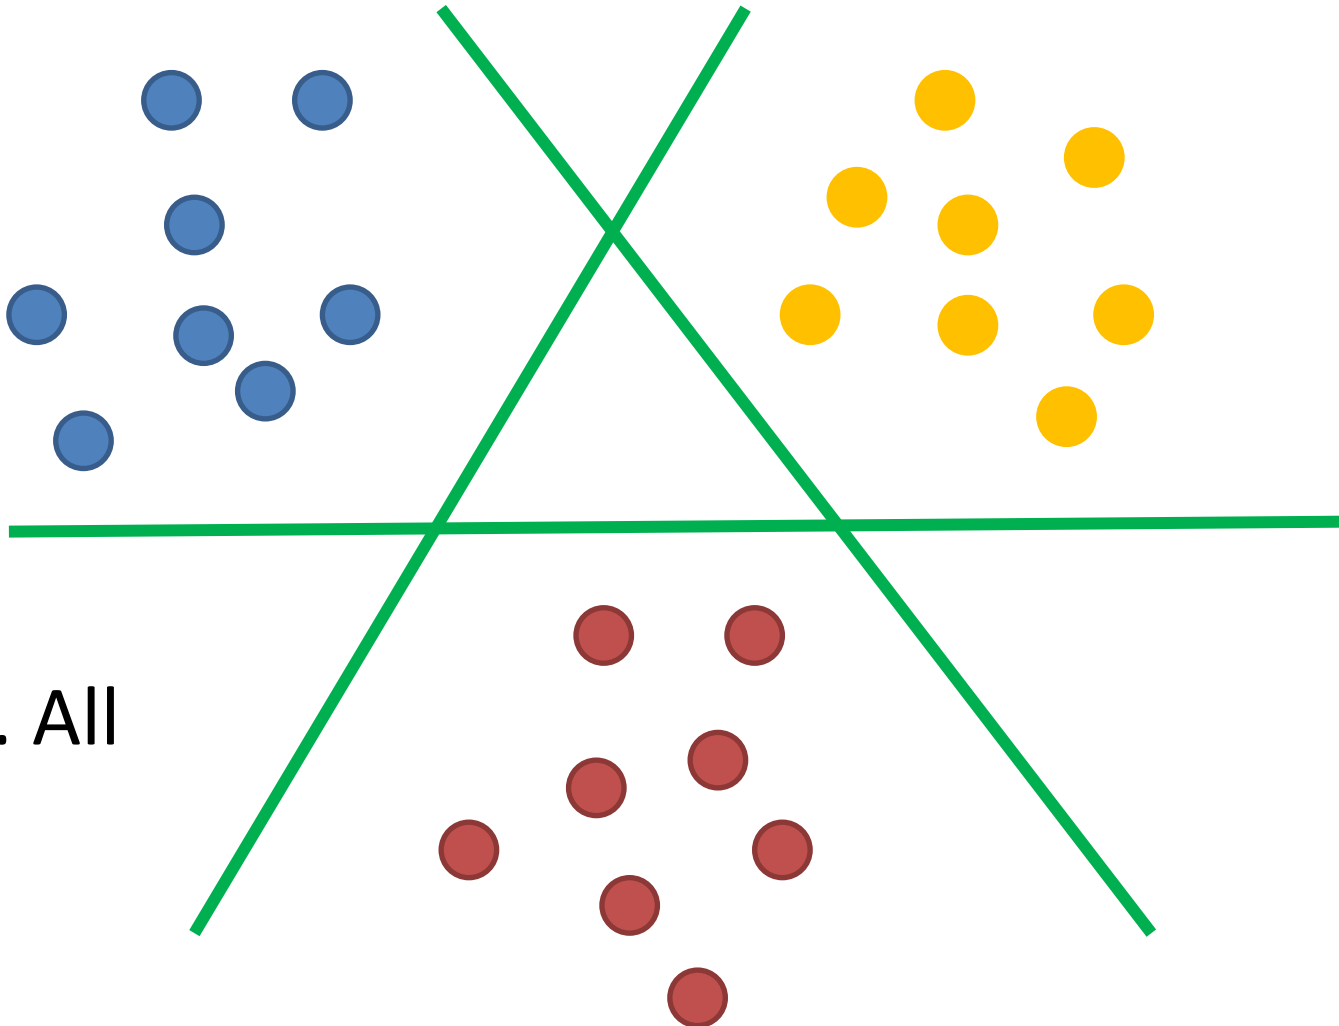
F-measure

- Weighted average of precision and recall

$$F_{\beta} = \frac{(\beta^2 + 1) \cdot P \cdot R}{\beta^2 \cdot P + R}$$

- Usual case: $\beta = 1$
- Increasing β allocates weight to recall

Multi Class

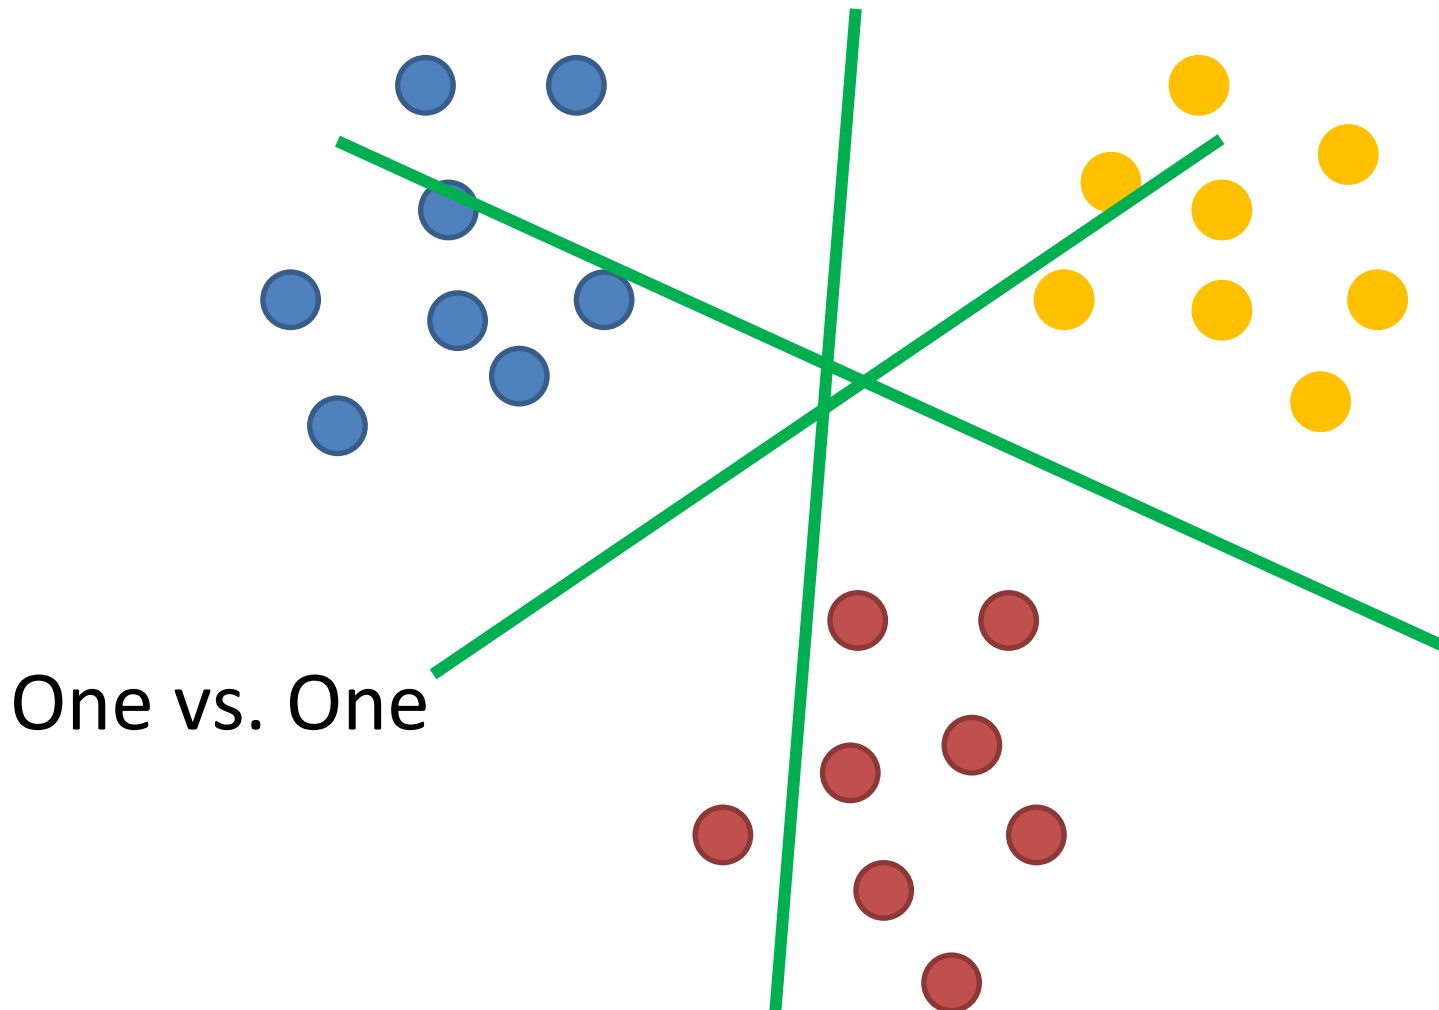


One vs. All

One vs All

- Train n classifier for n classes
- Take classification with greatest margin
- Slow training

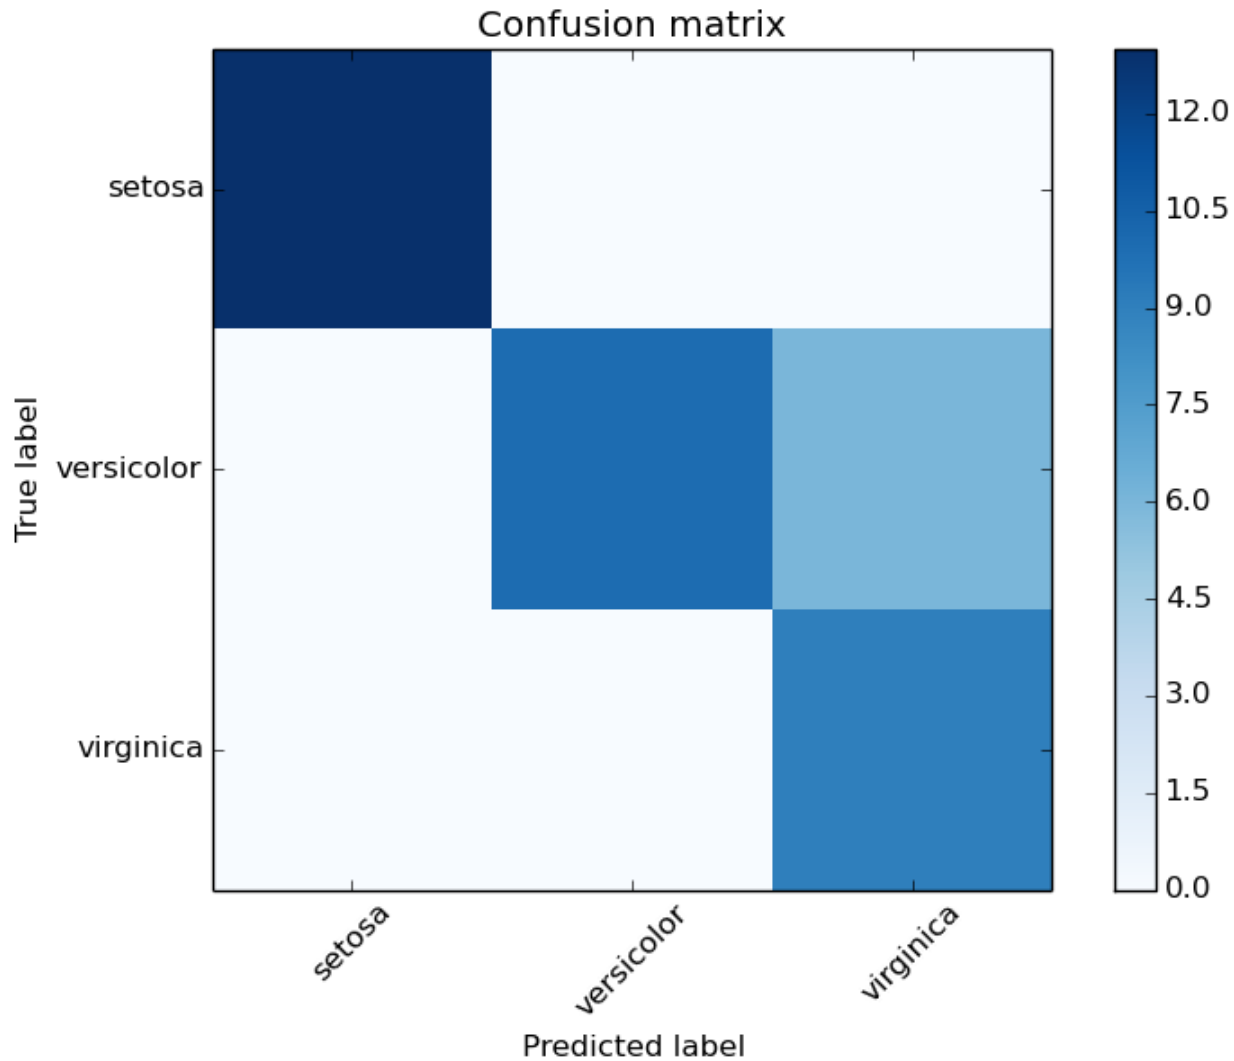
Multi Class



One vs One

- Train $n(n-1)/2$ classifiers
- Take majority vote
- Fast training

Confusion Matrix



Recap

- Perceptrons are great
- But really just a separating hyperplane
- So is SVM
- Kernels are neat
- Evaluation metrics are important